



Detecting Stealthy P2P Botnets Using Statistical Traffic Fingerprints

Junjie Zhang¹, Roberto Perdisci², Wenke Lee¹,
Unum Sarfraz¹, and Xiapu Luo³

¹Georgia Tech ²Univ. of Georgia

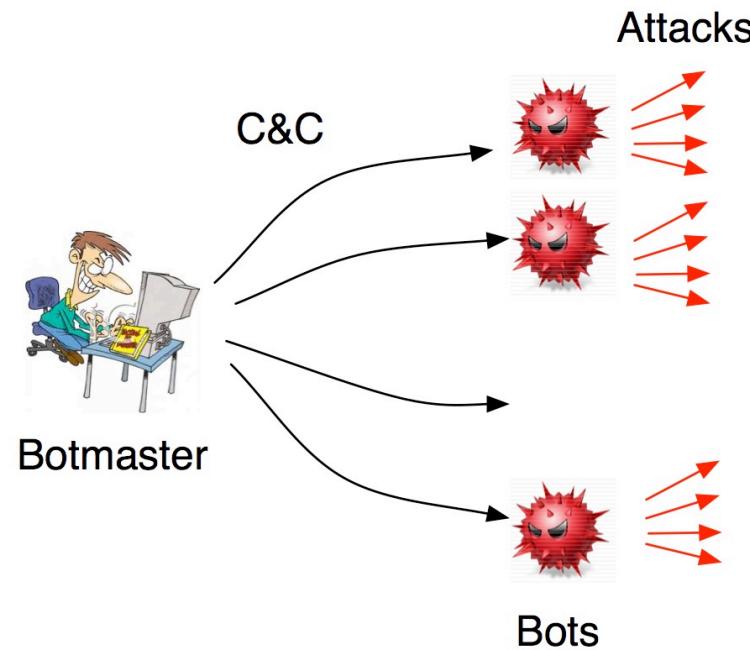
³Hong Kong Polytechnic Univ.

Outline

- Introduction
- System Design
- System Evaluation
- Conclusion

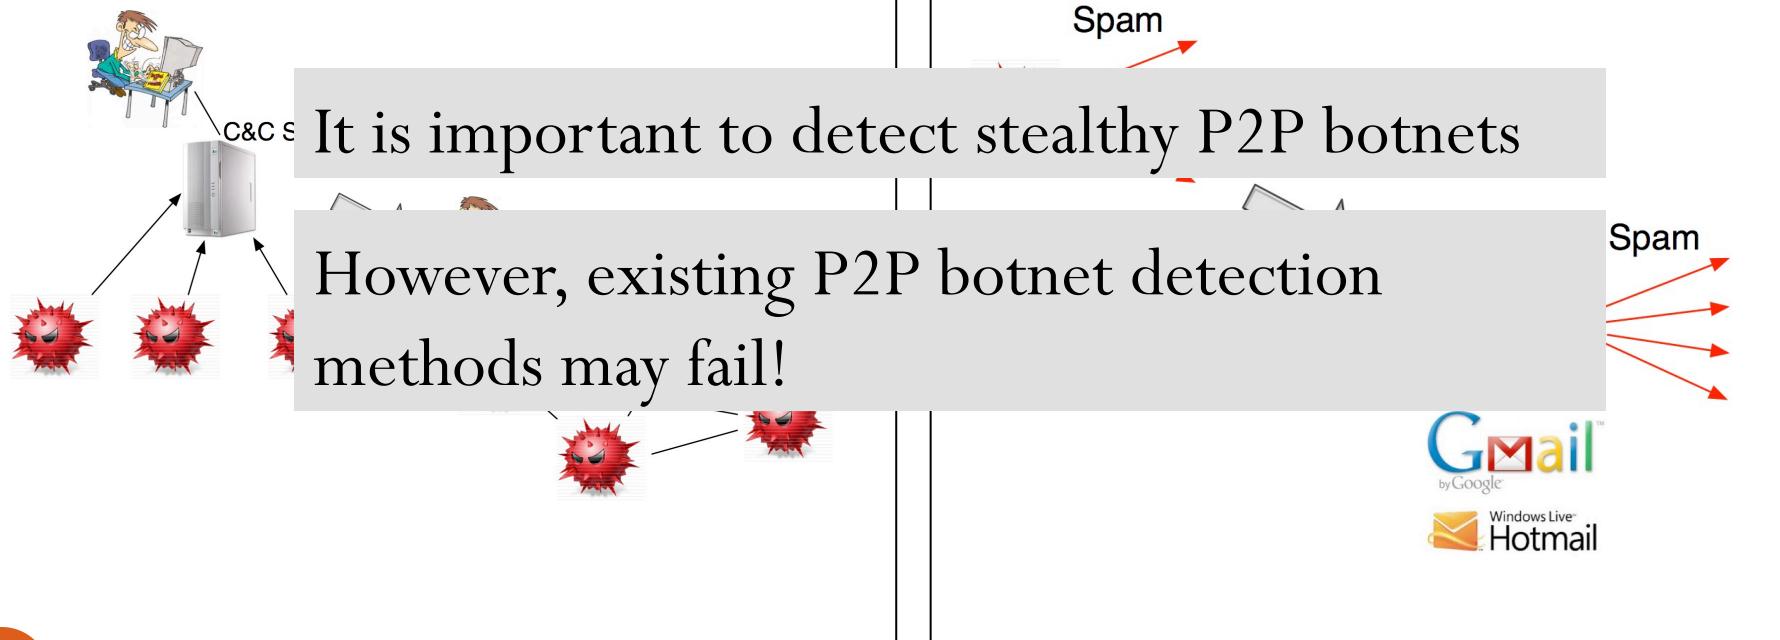
Introduction

- Botnet
 - A collection of bots controlled by a botmaster via a *command and control (C&C)* channel
 - Botnets serve as the infrastructures for a variety of attacks
 - Spam
 - Identity Theft
 - DDoS
 - ...



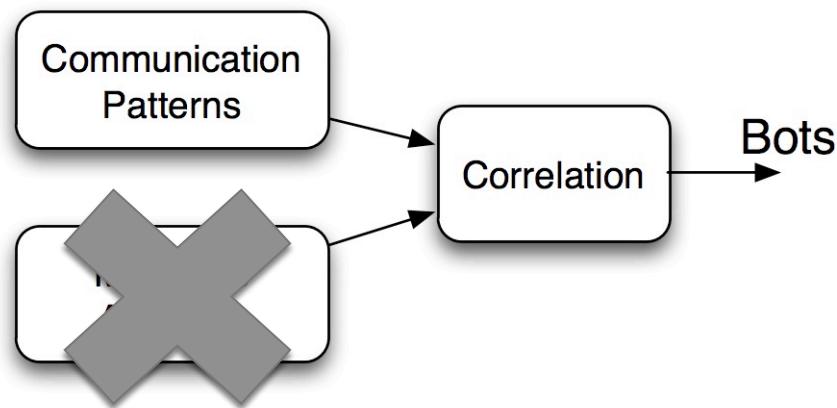
Introduction

- Botnets Evolve
- C&Cs
 - Centralized -> P2P
- Attacks
 - Noisy -> Stealthy



Introduction

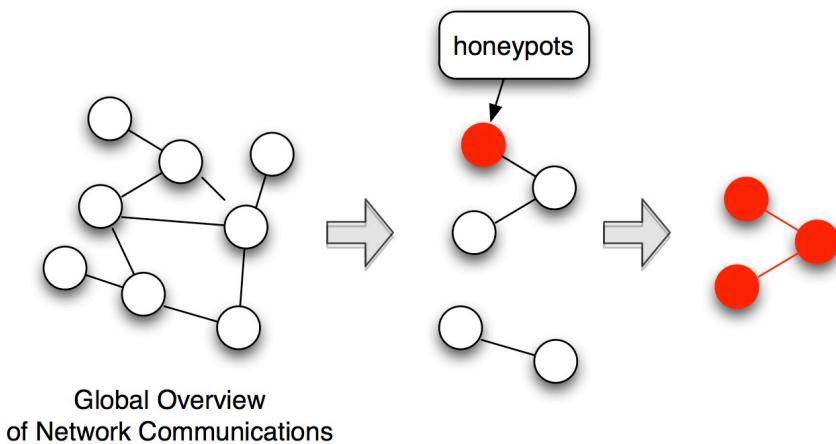
- Existing Method 1
- BotMiner [Gu et al Security 08]
 - Correlates similar activities in communication plane and attack plane



- However, stealthy attacks may not be observable, making the correlation ineffective!

Introduction

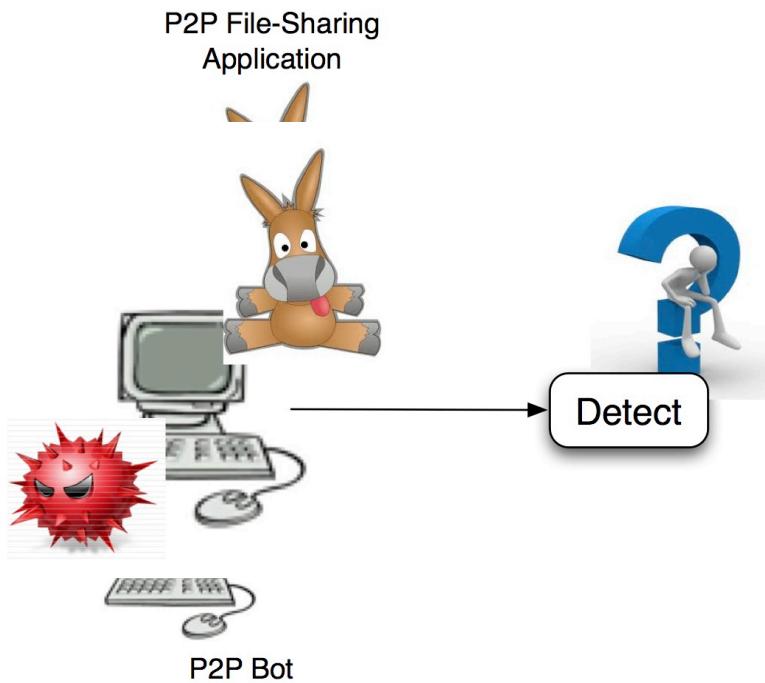
- Existing Method 2
- BotGrep [Nagaraja et al Security 10]
 - Uses graph analysis to identify P2P networks
 - Requires global overview of network communications
 - Requires bootstrap information (e.g., honeypots)



- However, it is challenging to get
 - Network global overview
 - A priori detection results
- Resulting in limited usage in practice!

Introduction

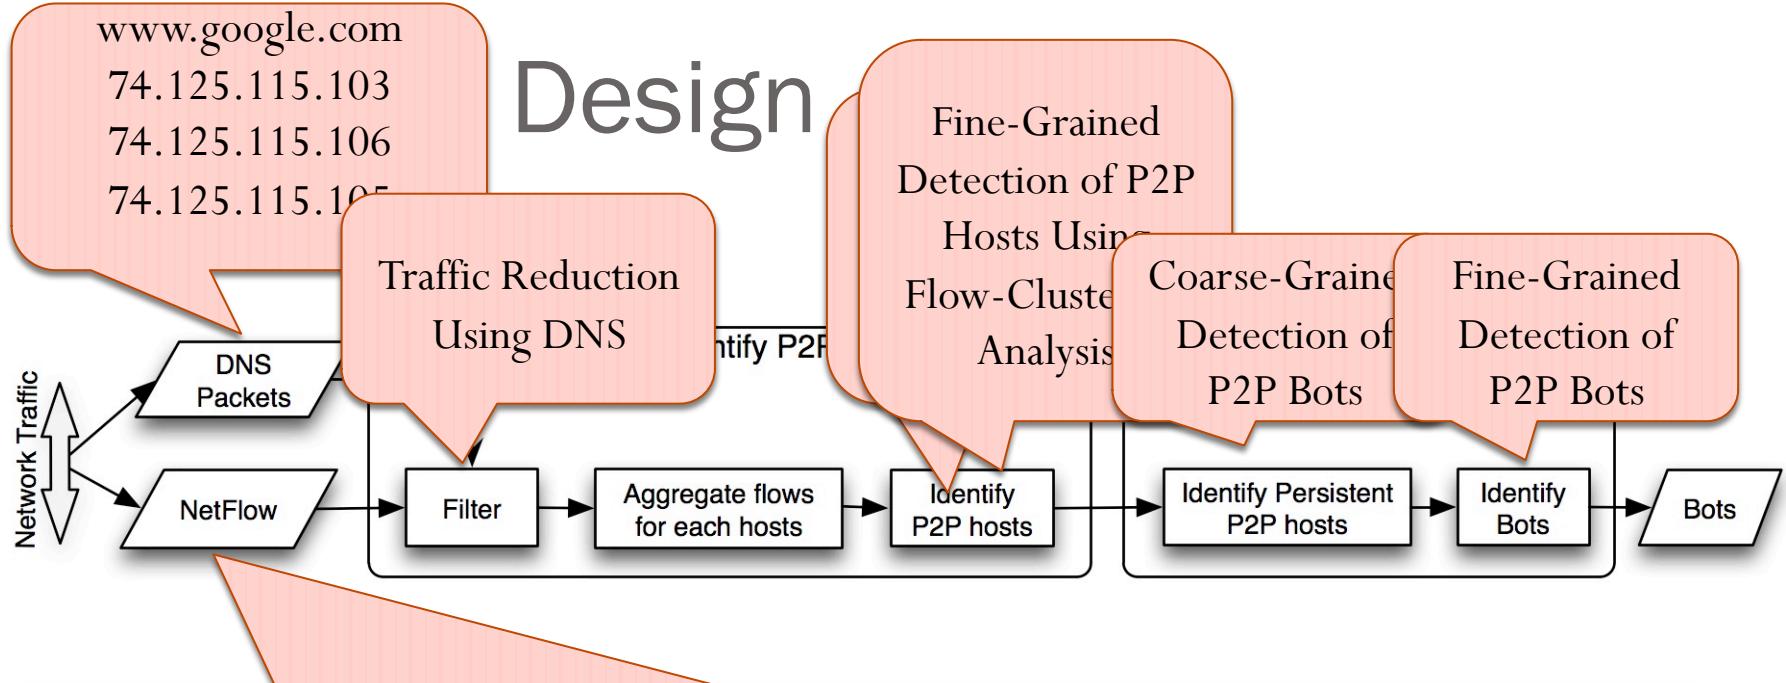
- Existing Method 3
- Telling P2P File-Sharing and Bots Apart [Yen et al ICDCS 10]
 - Differentiate P2P bots from P2P file-sharing clients



- However, the method may fail if bot-compromised host is running a P2P file-sharing application

Introduction

- We need a new method to detect stealthy P2P botnets!
- Design Targets:
 - Detect P2P bots without observing attack activities
 - Detect P2P bots even if the underlying hosts run legitimate P2P applications



Time	SrcIP	DstIP	SrcPort	DstPort	Proto	PktS	PktR	ByteS	ByteR	Flag
1305602966	192.168.1.5	74.125.115.103	2000	80	TCP	5	10	30	1000	SAF
1305602966	192.168.1.5	200.12.15.10	6000	1001	UDP	1	1	75	75	N/A
.....										

- Step 1.3. Fine-Grained Detection of P2P Hosts
- Phase II: Identify P2P Bots
 - Step 2.1. Coarse-Grained Detection of P2P Bots
 - Step 2.2 Fine-Grained Detection of P2P Bots

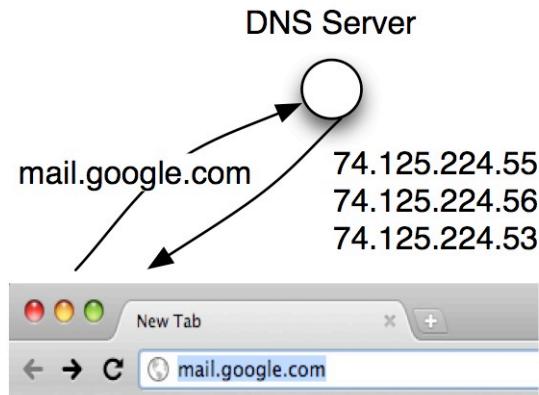
System Design

- Use real-world P2P traces to motivate the system design and parameter selection
 - 5 popular P2P applications
 - 24 hours

P2P Applications	Version	Protocol
Bittorrent	6.4	Bittorrent
Emule	0.49c	Kademlia
Limewire	5.4.8	Gnutella&Bittorrent
Skype	4.2	Skype
Ares	2.1.5	Gnutella&Bittorrent

System Design – Phase I: Identify P2P Hosts

- Step 1.1 Traffic Reduction
 - **Observation:** The peer IP addresses of P2P connections are obtained from the overlay routing tables instead of DNS responses.
 - **System Design:** Filter out the flows, whose destination IP addresses are resolved from DNS queries.



Trace	Percentage of flows associated with no-DNS DstIPs
T-Bittorrent	96.85%
T-Emule	99.99%
T-Limewire	99.97%
T-Skype	99.93%
T-Ares	99.99%

System Design – Phase I: Identify P2P Hosts

- Step 1.2 Coarse-Grained Detection of P2P Hosts
 - **Observation:** The significant node churn of P2P networks introduces a large number of failed connections.
 - **System Design:** Identify hosts that initiate more than θ_o failed connections for each day.

Trace	# of Failed Conns Per hour
T-Bittorrent	1602
T-Emule	318
T-Limewire	1278
T-Skype	81
T-Ares	489

Conservatively set $\theta_o = 10$

System Design – Phase I: Identify P2P Hosts

- Step 1.3 Fine-Grained Detection of P2P Hosts

- **Observation:**

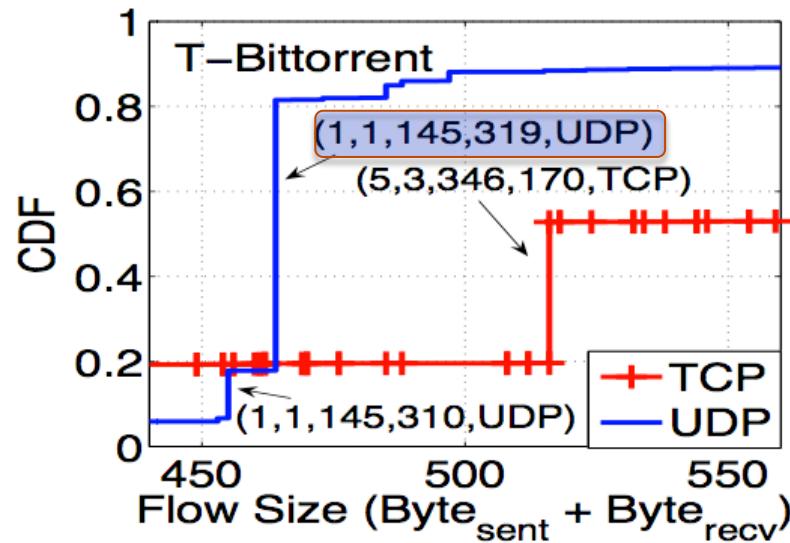
- A P2P application automatically generates a large number of control flows (e.g., ping/pong) *to a large number of peers*, as long as it is online.
 - The control flows used for the same purpose (e.g., ping/pong or peer discovery) *share similar size*.

- **System Design:**

- Aggregate flows with similar sizes to clusters respectively.
 - Define a cluster as *a fingerprint cluster*, if the number of unique BGP prefixes of the destination IP addresses for its flows is greater than θ_{BGP} ($=50$).
 - Claim a host as P2P host if it has at least one fingerprint cluster.

System Design – Phase I: Identify P2P Hosts

- An example of control flows

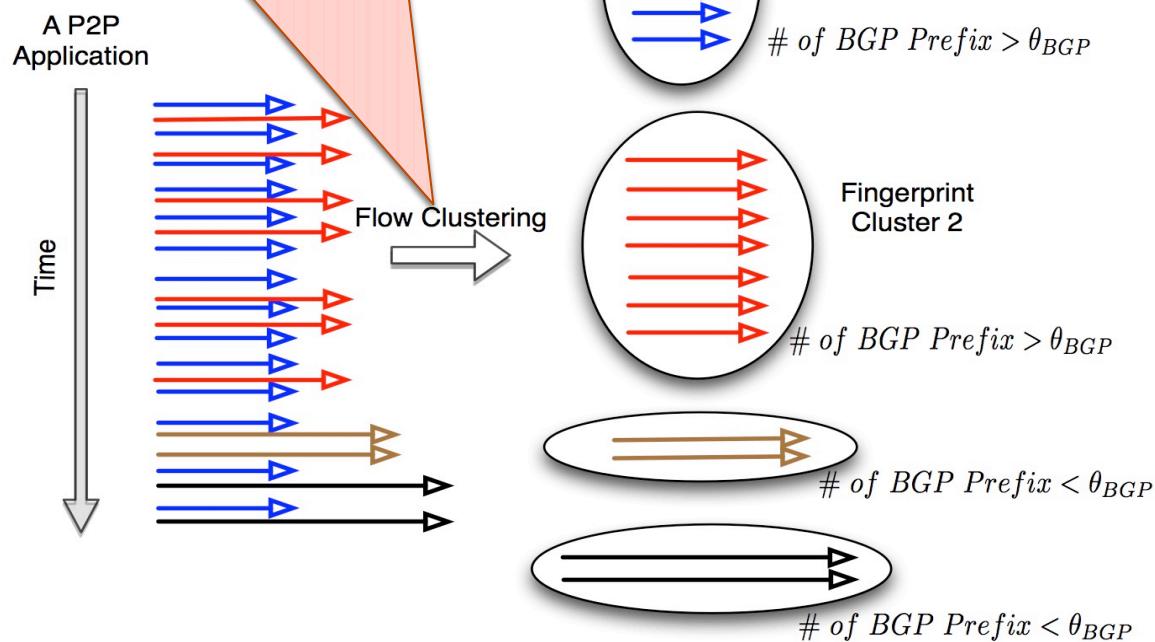


Fingerprints	flows	outgoing content	incoming content	description
1 1 145 319, UDP	1 2 ...	d1:ad2:id20:... find_node1:...:y1:qe	d1:rd2:id20:... nodes208:...:y1:re d1:rd2:id20:... nodes208:...:y1:re d1:rd2:id20:... nodes208:...:y1:re	peer discovery

System Design – Phase I: Identify P2P Hosts

- Step 1.3. Flow-Clustering analysis for identifying fingerprint

- Scalable clustering scheme (Birch + Hierarchical Clustering).
- Refer paper for details.



System Design – Phase II: Identify P2P Bots

- Step 2.1 Coarse-Grained Detection of P2P Bots
 - Motivation:
 - A bot usually is active as long as its underlying system is power on
 - A legitimate P2P application may only be used for a short period of time
 - System Design:
 - Identify P2P clients that are *persistently active* compared to its underlying host

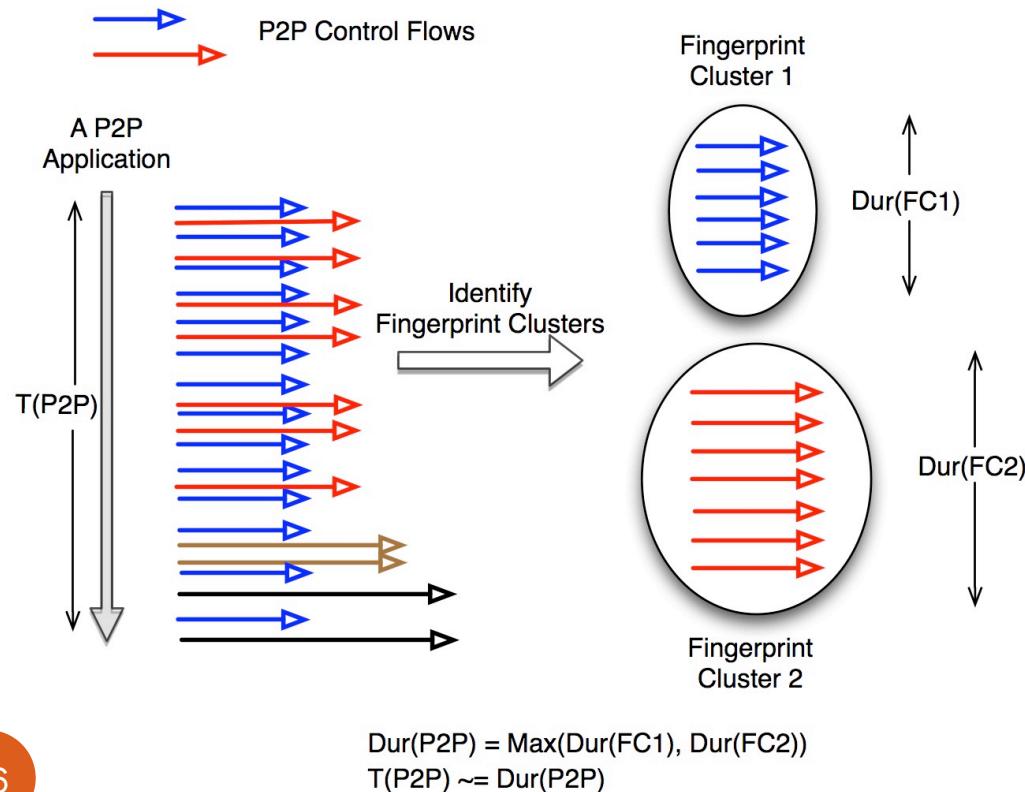
$$T(\text{P2P}) / T(\text{Host}) > 0.5$$

- Challenge: How can we accurately estimate the active time of a P2P application (i.e., $T(\text{P2P})$)?
 - Solution: Leverage obtained fingerprint clusters!

System Design – Phase II: Identify P2P Bots

- Estimate the active time of a P2P application ($T(P2P)$)

$$Dur(P2P) = \text{MAX}(Dur(FC_1), Dur(FC_2), \dots, Dur(FC_n))$$



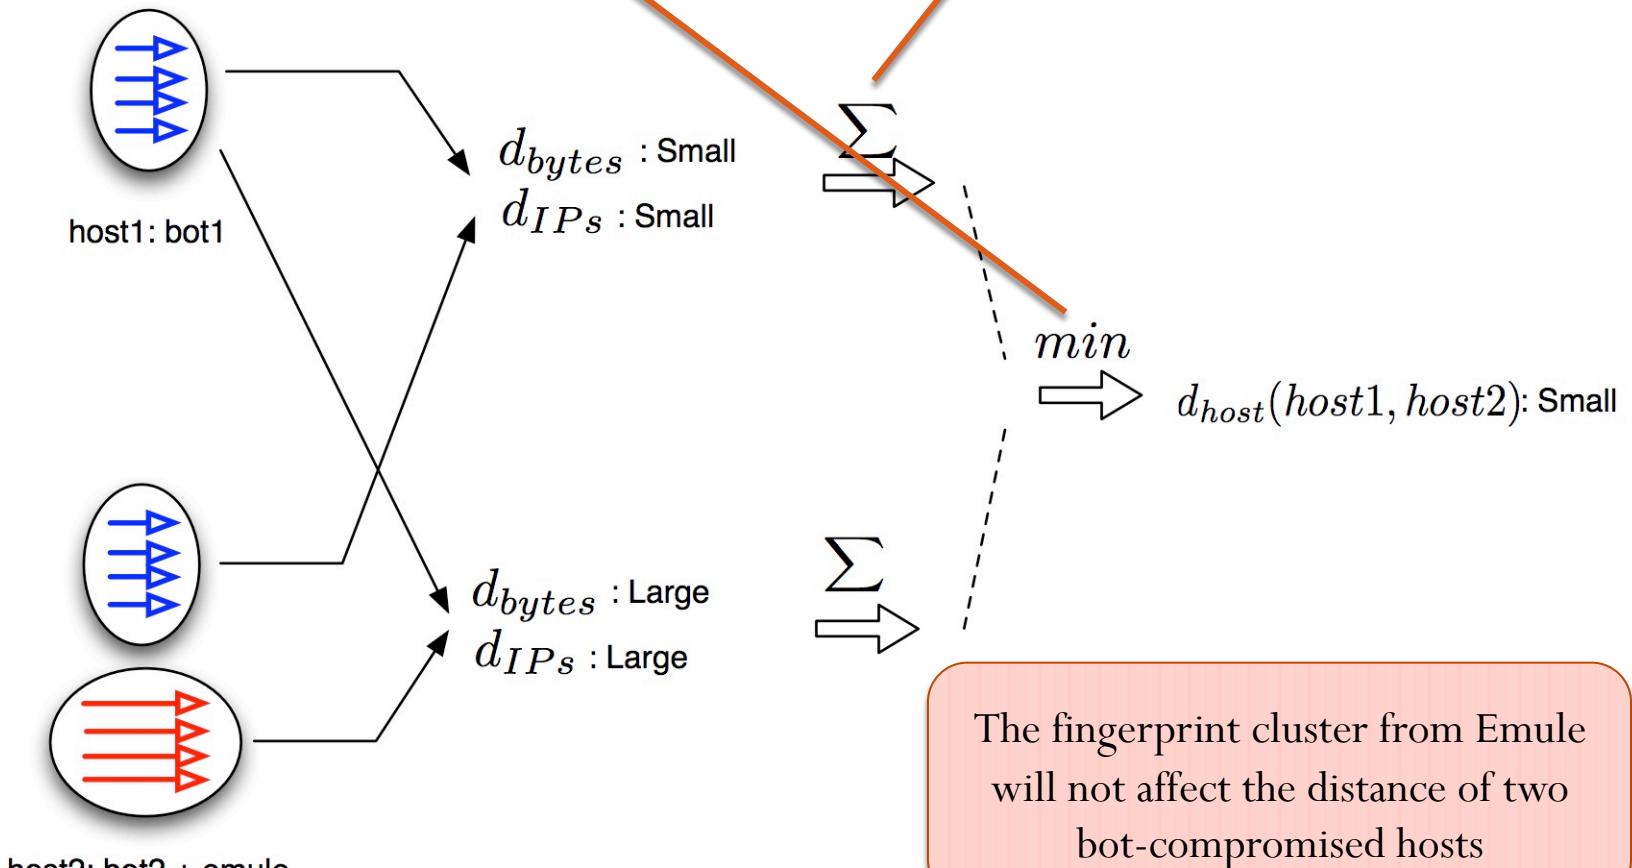
Trace	Actual P2P Active Time (T(P2P))	Estimated P2P Active Time (Dur(P2P))
T-Bittorrent	24hr	24hr
T-Emule	24hr	24hr
T-Limewire	24hr	24hr
T-Skype	24hr	24hr
T-Ares	24hr	24hr

- $T(P2P) \approx Dur(P2P)$.
- Dur(P2P) can accurately approximates T(P2P)!

System Design – Phase II: Identify P2P Bots

- Fine-Grained Detection of P2P Bots
 - Motivation
 - Same P2P applications, such as P2P bots of the same botnet, use the same P2P protocol, resulting in similar/same fingerprint clusters
 - P2P bots in the same botnet search for the same commands/contents. Therefore the search paths from bots tend to converge, resulting in a large overlap of peers.
 - Legitimate P2P clients are driven by different users, searching for different contents. Therefore the search paths for legitimate clients tend to diverge, resulting in a small overlap of peers.

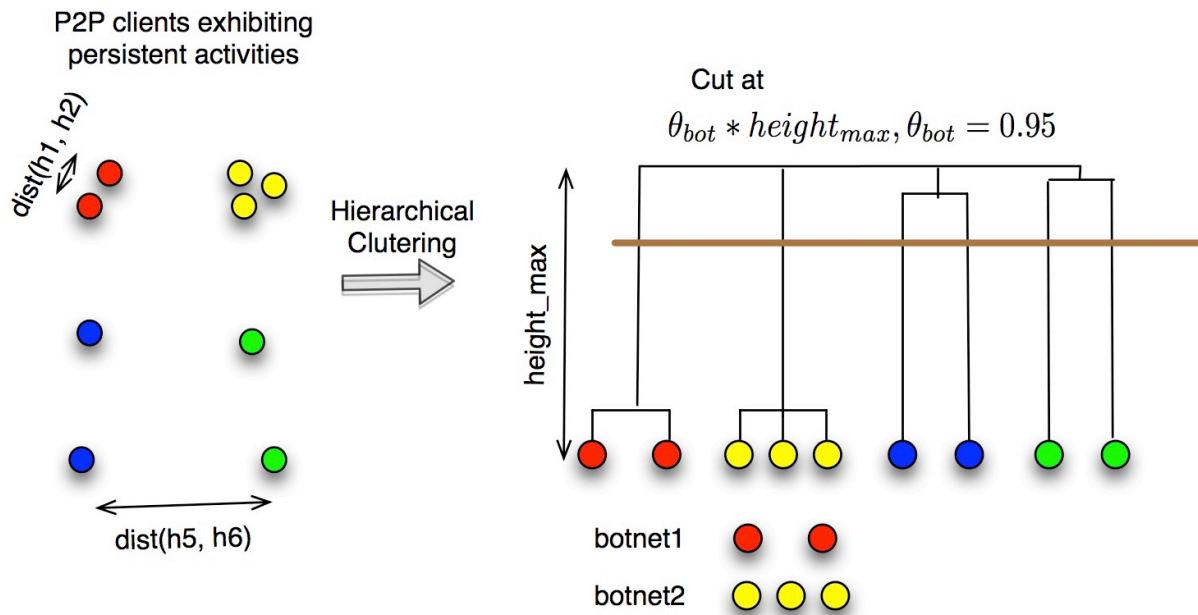
$$dist(h_a, h_b) = \min_{i,j} (\lambda * \frac{d_{bytes}(FC_i^{(a)}, FC_j^{(b)}) - \min_B}{\max_B - \min_B} + (1 - \lambda) * d_{IPs}(FC_i^{(a)}, FC_j^{(b)}))$$



bot-compromised system runs an Emule application

System Design – Phase II: Identify P2P Bots

- Fine-Grained Detection of P2P Bots
 - System Design: If two persistently active P2P clients have small distance, we identify them as bots



System Evaluation

- Experimental Setup

Data Sources	Description	P2P Clients/Traces
College networks	Passively collected for 24 hours; Netflows with 200 bytes payload between internal and external networks; DNS responses with IP addresses	3 BitTorrents (ground truth by signatures); 5 Skypes (potential)
5 P2P applications	Run each application <i>simultaneously</i> on two VMs for 24 (or 5) hours	10 traces totally
2 P2P botnets	Execute the malware binaries of Storm/Waledac for 24 hours	13 Storm bots; 3 Waledac bots

System Evaluation

- Experimental Design
 - Identifying and Profiling P2P applications
 - Detecting P2P bots
 - With being mixed with legitimate P2P data
 - Other situations are discussed in the paper including
 - Without being mixed with legitimate P2P data
 - Only two bots from each botnet
 - No bot (clean network)

System Evaluation

- Identifying and Profiling P2P Applications
 - DNS-based filter eliminates 2 / 3 hosts in the college networks!
 - Effectively identify and profile all the P2P applications!
 - 3 BitTorrent hosts in college networks
 - 5 potential Skypes in college networks
 - 10 P2P application instances we have run
 - 3 Waledac P2P bots
 - 13 Storm bots

Let us have a close look at Fingerprint Clusters!

System Evaluation

- Identifying and Profiling P2P Applications
 - Fingerprint cluster summaries for BitTorrent clients in college networks
 - Fingerprint cluster summary: average number of pkt_sent/recv, byte_sent/recv, and protocol for one fingerprint cluster

Fingerprints Ground truth		Fingerprints identified in the college network
(Pkt _s , Pkt _r , Byte _s , Byte _r , proto)		
T-BitTorrent	1 1 145 319, UDP	1 1 109 100, UDP
	1 1 109 100, UDP	1 1 146 340, UDP
T-BitTorrent-2	5 3 346 170, TCP	1 1 104 178, UDP
	1 1 145 310, UDP	1 1 319 145, UDP
T-BitTorrent-2	1 1 145.01 317.66, UDP	1 1 145 319, UDP
	1 1 109 100, UDP	1 1 145 319, UDP
T-BitTorrent-2	1 1 146 342, UDP	1 1 75 75, UDP
	5 3 346 170, TCP	1 1 65 65, UDP
T-BitTorrent-2	2 2 466 461, UDP	7 6 1118 1767, TCP

System Evaluation

- Identifying and Profiling P2P Applications
 - Fingerprint cluster summaries for Skype clients in college networks

Fingerprints Ground truth						
$(\overline{Pkt_s}, \overline{Pkt_r}, \overline{Bytes_s}, \overline{Bytes_r}, \text{proto})$						
Skype	<table border="1"><tr><td>1 1 74.58 60, UDP</td></tr><tr><td>1 1 78 60, UDP</td></tr><tr><td>1 1 75 60, UDP</td></tr><tr><td>1 1 76 60, UDP</td></tr><tr><td>1 1 79 60, UDP</td></tr></table>	1 1 74.58 60, UDP	1 1 78 60, UDP	1 1 75 60, UDP	1 1 76 60, UDP	1 1 79 60, UDP
1 1 74.58 60, UDP						
1 1 78 60, UDP						
1 1 75 60, UDP						
1 1 76 60, UDP						
1 1 79 60, UDP						

Fingerprints identified in
the college network
 $(\overline{Pkt_s}, \overline{Pkt_r}, \overline{Bytes_s}, \overline{Bytes_r}, \text{proto})$

Skype1@C	<table border="1"><tr><td>1 1 73 60, UDP</td></tr><tr><td>1 1 76 60, UDP</td></tr><tr><td>1 1 75 60, UDP</td></tr><tr><td>1 1 72 60, UDP</td></tr><tr><td>1 1 74 60, UDP</td></tr></table>	1 1 73 60, UDP	1 1 76 60, UDP	1 1 75 60, UDP	1 1 72 60, UDP	1 1 74 60, UDP
1 1 73 60, UDP						
1 1 76 60, UDP						
1 1 75 60, UDP						
1 1 72 60, UDP						
1 1 74 60, UDP						
Skype2@C	<table border="1"><tr><td>1 1 75 60, UDP</td></tr><tr><td>1 1 74 60, UDP</td></tr><tr><td>1 1 76 60, UDP</td></tr></table>	1 1 75 60, UDP	1 1 74 60, UDP	1 1 76 60, UDP		
1 1 75 60, UDP						
1 1 74 60, UDP						
1 1 76 60, UDP						
Skype3@C	<table border="1"><tr><td>1 1 72 60, UDP</td></tr><tr><td>1 1 74 60, UDP</td></tr><tr><td>1 1 79 60, UDP</td></tr><tr><td>1 1 76 60, UDP</td></tr></table>	1 1 72 60, UDP	1 1 74 60, UDP	1 1 79 60, UDP	1 1 76 60, UDP	
1 1 72 60, UDP						
1 1 74 60, UDP						
1 1 79 60, UDP						
1 1 76 60, UDP						
Skype4@C	<table border="1"><tr><td>1 1 73 60, UDP</td></tr></table>	1 1 73 60, UDP				
1 1 73 60, UDP						
Skype5@C	<table border="1"><tr><td>1 1 74 60, UDP</td></tr><tr><td>1 1 75 60, UDP</td></tr></table>	1 1 74 60, UDP	1 1 75 60, UDP			
1 1 74 60, UDP						
1 1 75 60, UDP						

System Evaluation

- Identifying and Profiling P2P Applications
 - Fingerprint cluster summaries for P2P bots

Fingerprints of P2P Bots			
(Pkt _s , Pkt _r , Byte _s , Byte _r , proto)		(Pkt _s , Pkt _r , Byte _s , Byte _r , proto)	
Storm1	2 2 94 554, UDP 2 2 94 1014, UDP 2 2 94 278, UDP ...	Storm2	2 2 94 554, UDP 2 2 94 1014, UDP 2 2 94 278, UDP ...
Waledac1	4 3 224 170, TCP 3 3 186 162, TCP 5 4 286 224, TCP ...	Waledac2	4 3 224 170, TCP 3 3 186 162, TCP 5 4 285 224, TCP ...

The fingerprint clusters can effectively identify and profile different P2P applications!

System Evaluation

- Detecting P2P bots
 - Overlaying legitimate P2P data and P2P botnet data to college network data.
 - half of bot-compromised hosts run legitimate P2P applications.

Bot	P2P App	Before Overlay (Bot)			After Overlay (Bot+P2PApp)		
		# of flows	# of DstIPs	avg flow size	# of flows	# of DstIPs	avg flow size
Waledac1	Emule1	341784	850	12829	452645	15338	55688
Waledac2	BT2@C	319119	760	11372	361135	1359	348708
Storm1	Limewire1	200237	6390	1342	429458	16635	1714
Storm2	BT3@C	275451	7319	1337	310667	8307	3381
Storm3	Bittorrent2	133955	5584	1344	432464	23261	172945
Storm4	Skype4@C	171471	7277	1280	199101	7520	1266
Storm5	Skype1	164917	6686	1328	214548	13137	1307
Storm6	Ares1	220459	6618	1307	238063	8543	6244

TABLE IX: Bot Traces Overlapped with P2P Application Traces

Bot traffic profiles are significantly skewed by the legitimate P2P data, which could make the existing detection approach such as [Yen et al ICDCS10] ineffective.

System Evaluation

- Detecting P2P Bots
 - Fingerprint clusters of P2P bots overlapped with legitimate P2P applications

Fingerprints of P2P Bots Mixed with P2P Apps

$(\overline{Pkt_s}, \overline{Pkt_r}, \overline{Byte_s}, \overline{Byte_r}, \text{proto})$

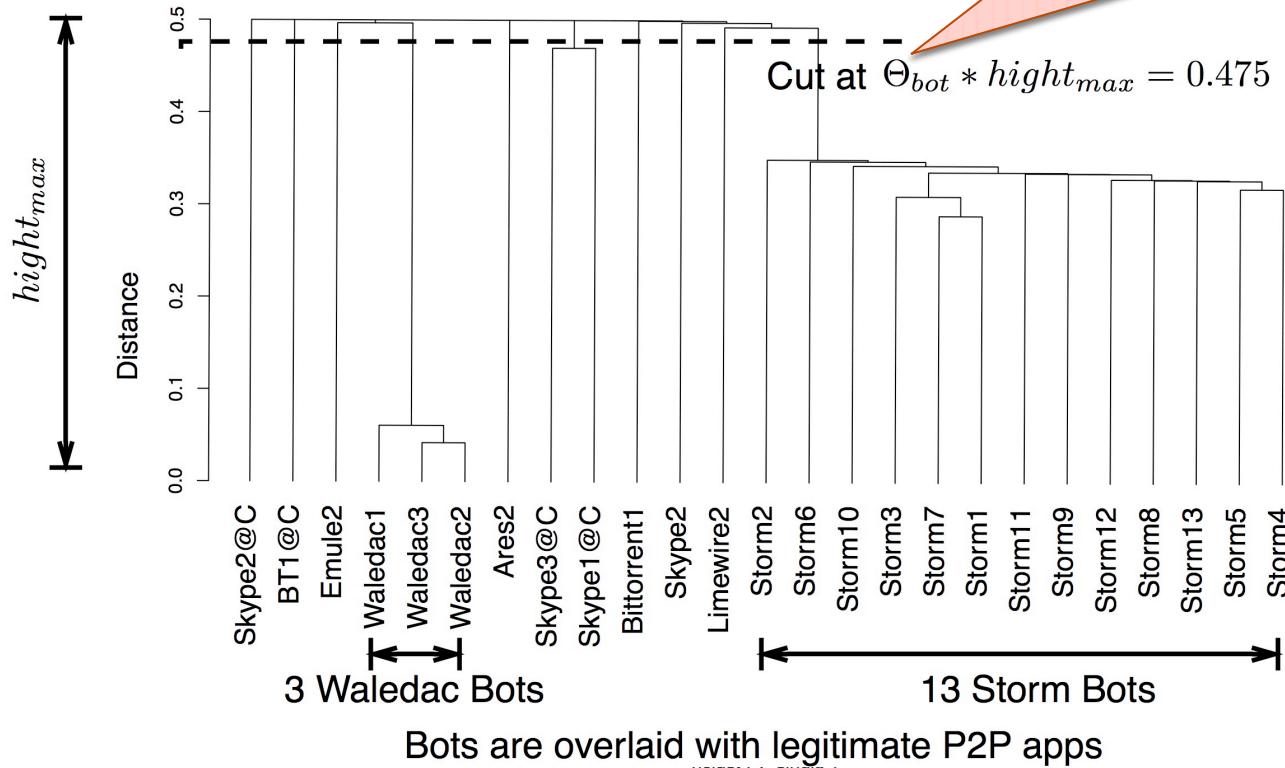
Waledac2+BT2@C	1 1 145 319, UDP (Bittorrent) 4 3 224 170, TCP (Waledac) 3 3 185 162, TCP (Waledac) 1 1 75 75, UDP (Bittorrent) ...
Storm4+Skype4@C	2 2 94 554, UDP (Storm) 2 2 94 1014, UDP (Storm) 1 1 73 60, UDP (Skype) ...

Fingerprint clusters can effectively profile P2P applications, even if the underlying host is running different applications.

System Evaluation

- Detecting P2P Bots

Refer paper for more evaluation of θ_{bot}

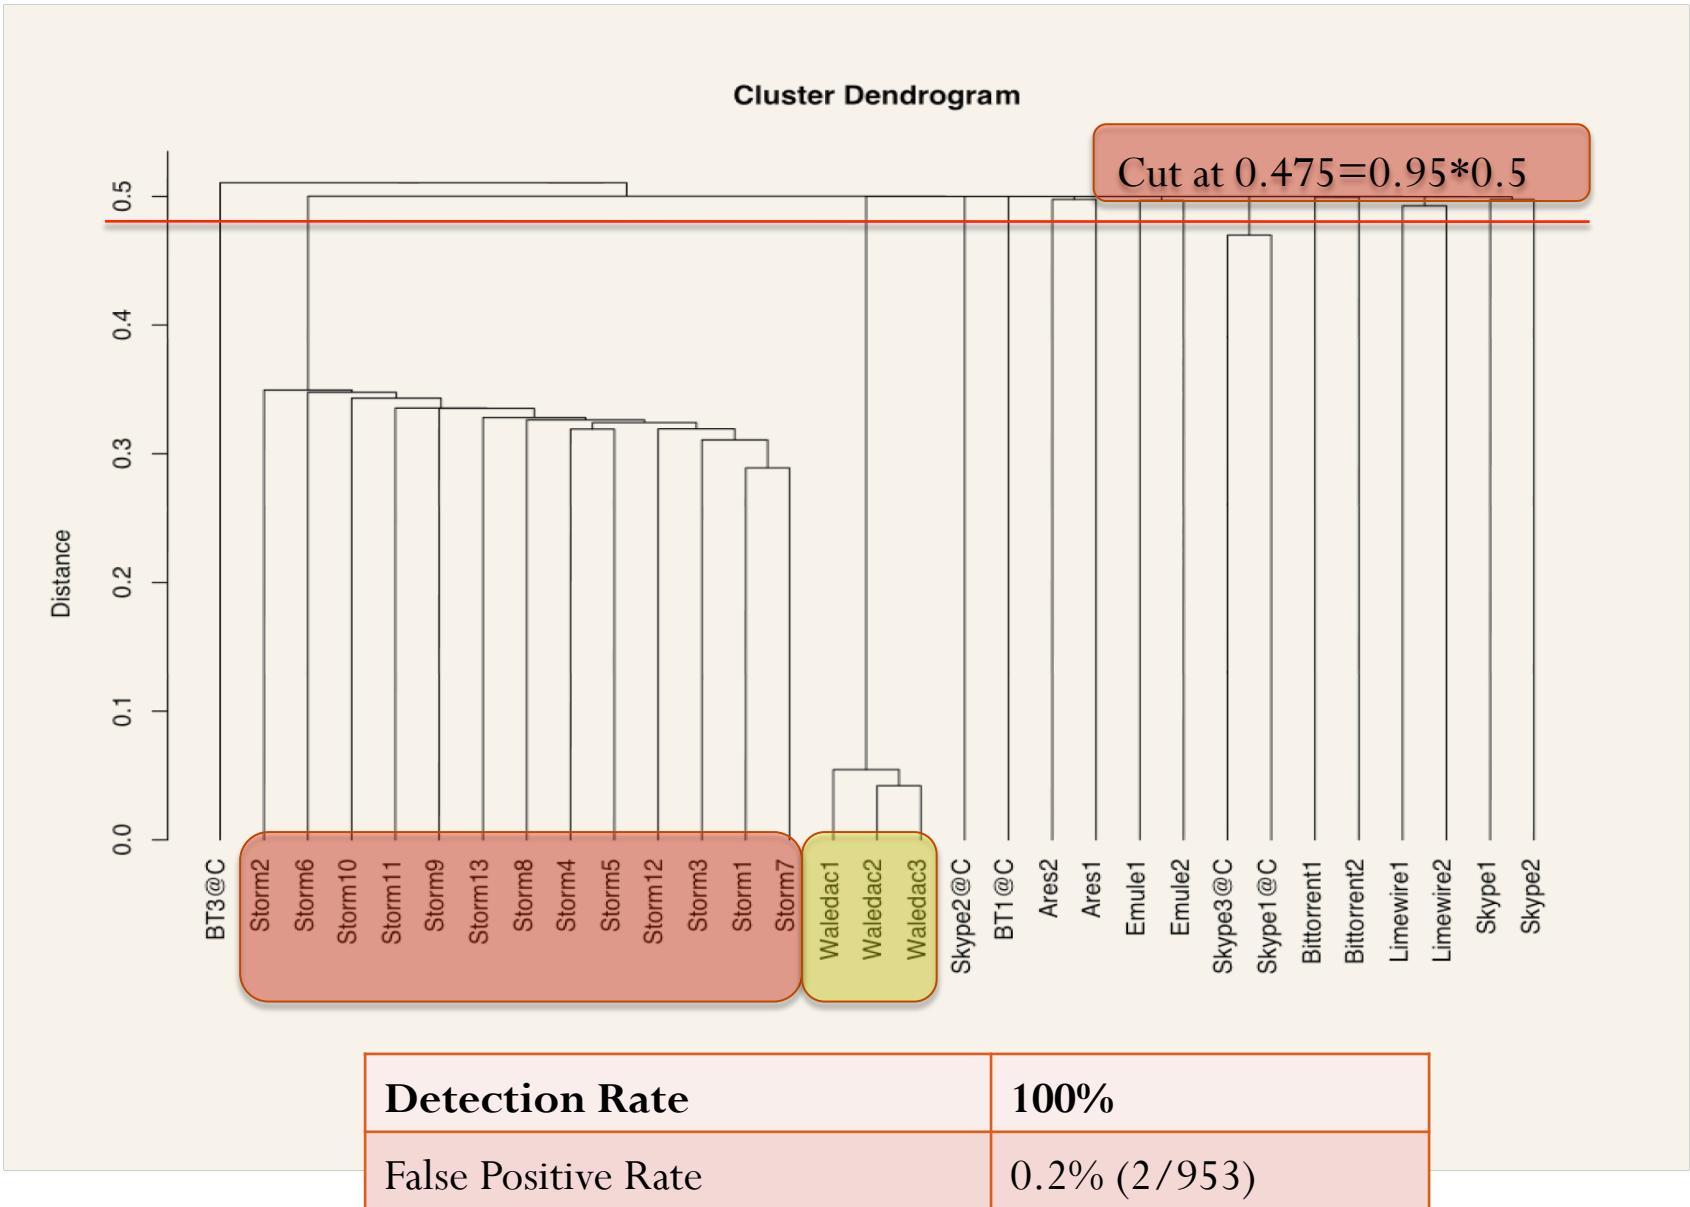


Conclusion

- A new P2P traffic detection algorithm
 - identify and *profile* P2P applications
- A novel P2P botnet detection system
 - Detect P2P bots without observing malicious activities
 - Detect P2P bots even if legitimate P2P applications are running on the underlying systems
 - High detection rate and low false positive rate

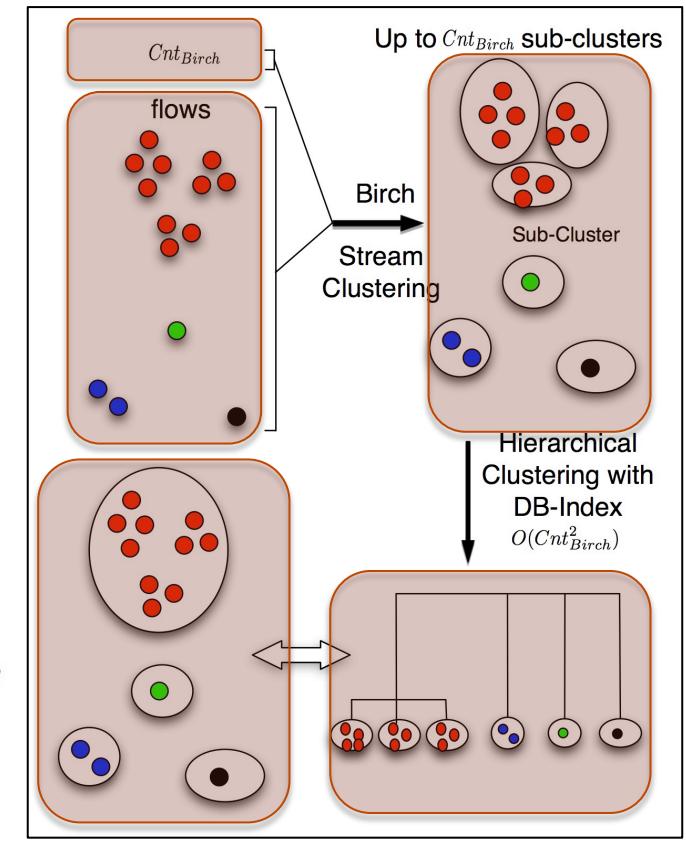
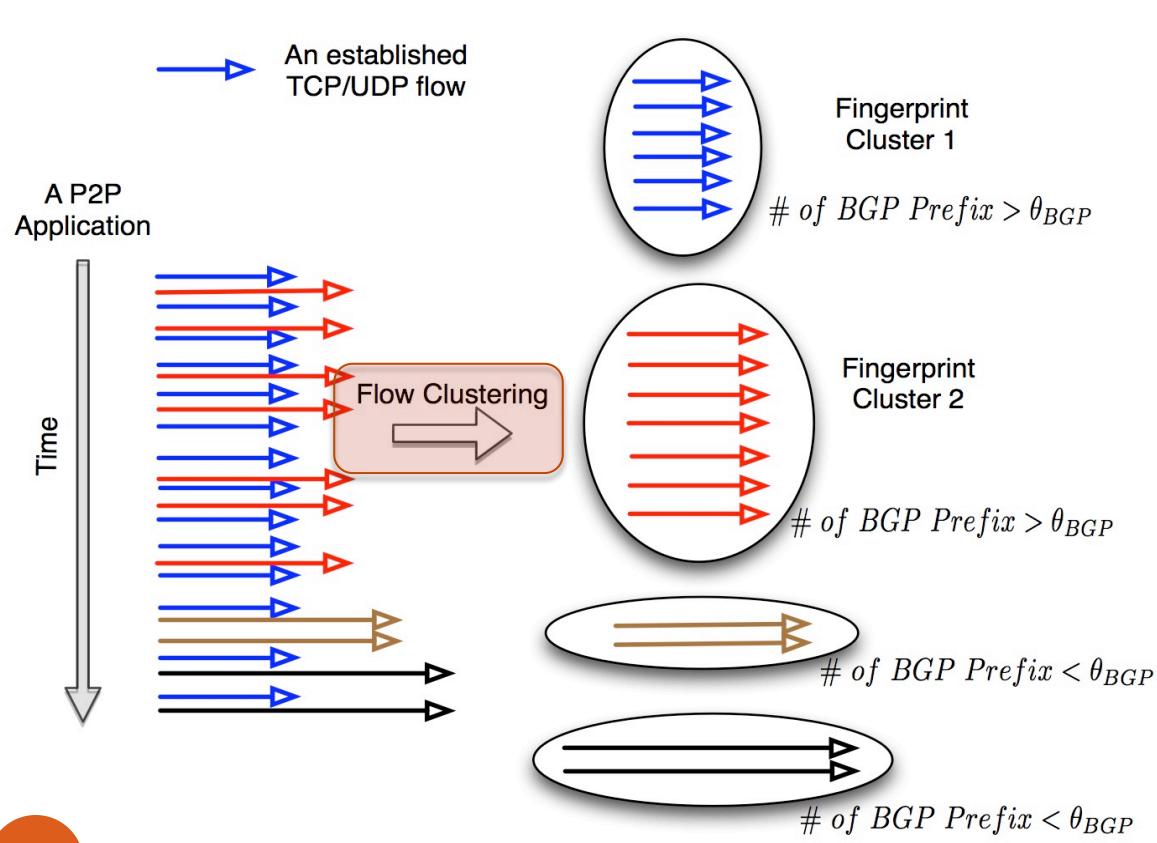
End

- Thanks!



System Design – Phase I: Identify P2P Hosts

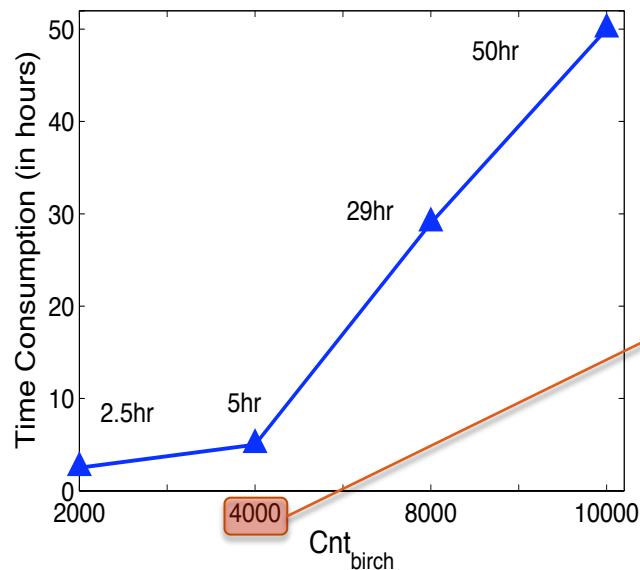
- Two steps of clustering analysis to generate fingerprint clusters.



$$Cnt_{Birch} = 4000$$

System Evaluation

- Experimental Results
 - Scalability Regarding Time Consumption



Our Experimental Configuration

Cnt_{birch}	-	0.1	0.3	0.5	Θ_{bot} 0.7	0.8	0.9	0.95
2000	DR FP	0 0	0 0	2/16 0	3/16 0	16/16 0	16/16 0	16/16 2/953
4000	DR FP	2/16 0	3/16 0	3/16 0	16/16 0	16/16 0	16/16 0	16/16 2/953
8000	DR FP	2/16 0	3/16 0	3/16 0	16/16 0	16/16 0	16/16 0	16/16 2/953
10000	DR FP	2/16 0	3/16 0	3/16 0	16/16 0	16/16 0	16/16 0	16/16 2/953

A small Cnt_{birch} increases system performance, but may decrease accuracy of fingerprint clusters.